

MELD Dataset Exploratory Data Analysis for Emotion Recognition

Sentiment and emotion analysis for customer service chatbot conversations

1. Import Required Libraries

```
In [ ]: #!/pip install wget
```

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
import zipfile
from collections import Counter
import warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: # Set style for better visualizations
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette("husl")
```

2. Download and Extract MELD Dataset

```
In [ ]: '''
if not os.path.exists('data'):
    os.makedirs('data')
# Download MELD dataset if not already present
meld_url = "http://web.eecs.umich.edu/~mihalcea/downloads/MELD.Raw.tar.gz"
output_path = "data/MELD.Raw.tar.gz"

if not os.path.exists(output_path):
    print("Downloading MELD dataset...")
    wget.download(meld_url, output_path)
    print("\nDownload completed!")
else:
    print("Dataset already downloaded.")

# Extract the dataset
import tarfile
if not os.path.exists('data/MELD.Raw'):
    print("Extracting dataset...")
    with tarfile.open(output_path, 'r:gz') as tar:
        tar.extractall('data/')
```

```
    print("Extraction completed!")
    ...
```

```
Out[ ]: '\nif not os.path.exists(\'data\'):\n    os.makedirs(\'data\')\n# Download ME  
LD dataset if not already present\nmeld_url = "http://web.eecs.umich.edu/~mih  
alcea/downloads/MELD.Raw.tar.gz"\noutput_path = "data/MELD.Raw.tar.gz"\n\nif  
not os.path.exists(output_path):\n    print("Downloading MELD dataset...")\n    wget.download(meld_url, output_path)\n    print("\nDownload completed!")\nelse:\n    print("Dataset already downloaded.")\n\n# Extract the dataset\nimport  
tarfile\nif not os.path.exists(\'data/MELD.Raw\'):\n    print("Extracting dat  
aset...")\n    with tarfile.open(output_path, \'r:gz\') as tar:\n        tar.  
extractall(\'data/\')\n    print("Extraction completed!")\n'
```

3. Load Dataset Files

```
In [ ]: # Define paths to the CSV files
# Use raw GitHub URLs for direct file access
train_path = 'https://raw.githubusercontent.com/declare-lab/MELD/master/data/M
dev_path = 'https://raw.githubusercontent.com/declare-lab/MELD/master/data/MEI
test_path = 'https://raw.githubusercontent.com/declare-lab/MELD/master/data/MEI

# Load the datasets
print("Loading datasets...")
train_df = pd.read_csv(train_path)
dev_df = pd.read_csv(dev_path)
test_df = pd.read_csv(test_path)
```

```
print(f"Train set loaded: {train_df.shape}")
print(f"Dev set loaded: {dev_df.shape}")
print(f"Test set loaded: {test_df.shape}")
```

Loading datasets...

Train set loaded: (9989, 11)

Dev set loaded: (1109, 11)

Test set loaded: (2610, 11)

4. Initial Data Exploration

```
In [ ]: # Display basic information about the training dataset
print("=== Training Dataset Info ===")
print(train_df.info())
print("\n=== First 5 rows of training data ===")
print(train_df.head())

# Check column names
print("\n=== Column Names ===")
print(train_df.columns.tolist())
```

=== Training Dataset Info ===

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 9989 entries, 0 to 9988

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Sr No.	9989 non-null	int64
1	Utterance	9989 non-null	object
2	Speaker	9989 non-null	object
3	Emotion	9989 non-null	object
4	Sentiment	9989 non-null	object
5	Dialogue_ID	9989 non-null	int64
6	Utterance_ID	9989 non-null	int64
7	Season	9989 non-null	int64
8	Episode	9989 non-null	int64
9	StartTime	9989 non-null	object
10	EndTime	9989 non-null	object

dtypes: int64(5), object(6)

memory usage: 858.6+ KB

None

=== First 5 rows of training data ===

	Sr No.	Utterance	Speaker
\			
0	1	also I was the point person on my company's tr...	Chandler
1	2	You must've had your hands full.	The Interviewer
2	3	That I did. That I did.	Chandler
3	4	So let's talk a little bit about your duties.	The Interviewer
4	5	My duties? All right.	Chandler

	Emotion	Sentiment	Dialogue_ID	Utterance_ID	Season	Episode	\
0	neutral	neutral	0	0	8	21	
1	neutral	neutral	0	1	8	21	
2	neutral	neutral	0	2	8	21	
3	neutral	neutral	0	3	8	21	
4	surprise	positive	0	4	8	21	

	StartTime	EndTime
0	00:16:16,059	00:16:21,731
1	00:16:21,940	00:16:23,442
2	00:16:23,442	00:16:26,389
3	00:16:26,820	00:16:29,572
4	00:16:34,452	00:16:40,917

=== Column Names ===

```
['Sr No.', 'Utterance', 'Speaker', 'Emotion', 'Sentiment', 'Dialogue_ID', 'Utterance_ID', 'Season', 'Episode', 'StartTime', 'EndTime']
```

```
In [ ]: # Display basic information about the training dataset
print("=== Training Dataset Info ===")
print(train_df.info())
print("\n=== First 5 rows of training data ===")
print(train_df.head())

# Check column names
print("\n=== Column Names ===")
print(train_df.columns.tolist())
```

=== Training Dataset Info ===

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 9989 entries, 0 to 9988

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Sr No.	9989 non-null	int64
1	Utterance	9989 non-null	object
2	Speaker	9989 non-null	object
3	Emotion	9989 non-null	object
4	Sentiment	9989 non-null	object
5	Dialogue_ID	9989 non-null	int64
6	Utterance_ID	9989 non-null	int64
7	Season	9989 non-null	int64
8	Episode	9989 non-null	int64
9	StartTime	9989 non-null	object
10	EndTime	9989 non-null	object

dtypes: int64(5), object(6)

memory usage: 858.6+ KB

None

=== First 5 rows of training data ===

	Sr No.	Utterance	Speaker
\			
0	1	also I was the point person on my company's tr...	Chandler
1	2	You must've had your hands full.	The Interviewer
2	3	That I did. That I did.	Chandler
3	4	So let's talk a little bit about your duties.	The Interviewer
4	5	My duties? All right.	Chandler

	Emotion	Sentiment	Dialogue_ID	Utterance_ID	Season	Episode	\
0	neutral	neutral	0	0	8	21	
1	neutral	neutral	0	1	8	21	
2	neutral	neutral	0	2	8	21	
3	neutral	neutral	0	3	8	21	
4	surprise	positive	0	4	8	21	

	StartTime	EndTime
0	00:16:16,059	00:16:21,731
1	00:16:21,940	00:16:23,442
2	00:16:23,442	00:16:26,389
3	00:16:26,820	00:16:29,572
4	00:16:34,452	00:16:40,917

=== Column Names ===

```
['Sr No.', 'Utterance', 'Speaker', 'Emotion', 'Sentiment', 'Dialogue_ID', 'Utterance_ID', 'Season', 'Episode', 'StartTime', 'EndTime']
```

5. Dataset Statistics

```
In [ ]: # Combined dataset statistics
total_utterances = len(train_df) + len(dev_df) + len(test_df)
print(f"Total utterances in dataset: {total_utterances}")
print(f"Training set: {len(train_df)} ({len(train_df)/total_utterances*100:.1f}%")
print(f"Development set: {len(dev_df)} ({len(dev_df)/total_utterances*100:.1f}%")
print(f"Test set: {len(test_df)} ({len(test_df)/total_utterances*100:.1f}%")
```



```
# Check for missing values
print("\n=== Missing Values in Training Set ===")
print(train_df.isnull().sum())
```

Total utterances in dataset: 13708

Training set: 9989 (72.9%)

Development set: 1109 (8.1%)

Test set: 2610 (19.0%)

=== Missing Values in Training Set ===

Sr No.	0
Utterance	0
Speaker	0
Emotion	0
Sentiment	0
Dialogue_ID	0
Utterance_ID	0
Season	0
Episode	0
StartTime	0
EndTime	0
dtype:	int64

6. Emotion Distribution Analysis

```
In [ ]: # Analyze emotion distribution across datasets
def plot_emotion_distribution(train_df, dev_df, test_df):
```

```
fig, axes = plt.subplots(1, 3, figsize=(18, 5))


# Training set
emotion_counts_train = train_df['Emotion'].value_counts()
axes[0].bar(emotion_counts_train.index, emotion_counts_train.values)
axes[0].set_title('Emotion Distribution - Training Set', fontsize=14)
axes[0].set_xlabel('Emotion')
axes[0].set_ylabel('Count')
axes[0].tick_params(axis='x', rotation=45)

# Dev set
emotion_counts_dev = dev_df['Emotion'].value_counts()
axes[1].bar(emotion_counts_dev.index, emotion_counts_dev.values)
axes[1].set_title('Emotion Distribution - Development Set', fontsize=14)
axes[1].set_xlabel('Emotion')
axes[1].set_ylabel('Count')
axes[1].tick_params(axis='x', rotation=45)

# Test set
emotion_counts_test = test_df['Emotion'].value_counts()
axes[2].bar(emotion_counts_test.index, emotion_counts_test.values)
axes[2].set_title('Emotion Distribution - Test Set', fontsize=14)
axes[2].set_xlabel('Emotion')
axes[2].set_ylabel('Count')
axes[2].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()
```

```
plot_emotion_distribution(train_df, dev_df, test_df)
```



```
In [ ]: # Overall emotion distribution
all_emotions = pd.concat([train_df['Emotion'], dev_df['Emotion'], test_df['Emo
emotion_dist = all_emotions.value_counts()
print("\n=== Overall Emotion Distribution ===")
print(emotion_dist)
print(f"\nTotal unique emotions: {len(emotion_dist)}")
```

=== Overall Emotion Distribution ===

Emotion

neutral 6436

joy 2308

surprise 1636

anger 1607

sadness 1002

disgust 361

fear 358

Name: count, dtype: int64

Total unique emotions: 7

7. Sentiment Distribution Analysis

```
In [ ]: # Analyze sentiment distribution
def plot_sentiment_distribution(train_df, dev_df, test_df):
    fig, axes = plt.subplots(1, 3, figsize=(18, 5))


    datasets = [('Training', train_df), ('Development', dev_df), ('Test', test_df)]

    for idx, (name, df) in enumerate(datasets):
        sentiment_counts = df['Sentiment'].value_counts()
        axes[idx].pie(sentiment_counts.values, labels=sentiment_counts.index,
                      axes[idx].set_title(f'Sentiment Distribution - {name} Set', fontsize=12))

    plt.tight_layout()
```

```
plt.show()
```

```
plot_sentiment_distribution(train_df, dev_df, test_df)
```



```
In [ ]: # Overall sentiment distribution
all_sentiments = pd.concat([train_df['Sentiment'], dev_df['Sentiment'], test_df['Sentiment']])
sentiment_dist = all_sentiments.value_counts()
print("\n=== Overall Sentiment Distribution ===")
print(sentiment_dist)
```

```
=== Overall Sentiment Distribution ===
```

```
Sentiment
```

```
neutral      6436
```

```
negative     4184
```

```
positive     3088
```

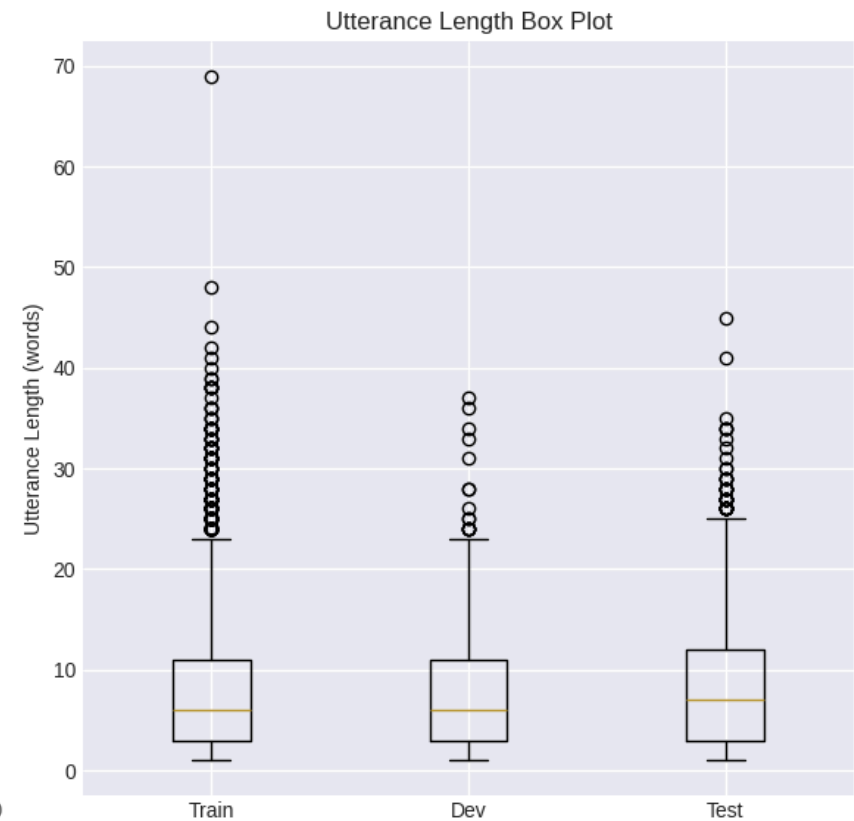

```
Name: count, dtype: int64
```

8. Text Length Analysis

```
In [ ]: # Analyze utterance lengths
train_df['utterance_length'] = train_df['Utterance'].apply(lambda x: len(str(x)))
dev_df['utterance_length'] = dev_df['Utterance'].apply(lambda x: len(str(x).split()))
test_df['utterance_length'] = test_df['Utterance'].apply(lambda x: len(str(x).split()))
```

```
In [ ]: # Plot utterance length distribution
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.hist(train_df['utterance_length'], bins=50, alpha=0.7, label='Train')
plt.hist(dev_df['utterance_length'], bins=50, alpha=0.7, label='Dev')
plt.hist(test_df['utterance_length'], bins=50, alpha=0.7, label='Test')
plt.xlabel('Utterance Length (words)')
plt.ylabel('Frequency')
plt.title('Distribution of Utterance Lengths')
plt.legend()
plt.xlim(0, 100)

plt.subplot(1, 2, 2)
plt.boxplot([train_df['utterance_length'], dev_df['utterance_length'], test_df['utterance_length']],
            labels=['Train', 'Dev', 'Test'])
plt.ylabel('Utterance Length (words)')
plt.title('Utterance Length Box Plot')
plt.tight_layout()
plt.show()
```



```
In [ ]: # Statistics
print("=== Utterance Length Statistics ===")
print(f"Training set - Mean: {train_df['utterance_length'].mean():.2f}, Median: {train_df['utterance_length'].median():.2f}, Max: {train_df['utterance_length'].max():.2f}")
print(f"Dev set - Mean: {dev_df['utterance_length'].mean():.2f}, Median: {dev_df['utterance_length'].median():.2f}, Max: {dev_df['utterance_length'].max():.2f}")
print(f"Test set - Mean: {test_df['utterance_length'].mean():.2f}, Median: {test_df['utterance_length'].median():.2f}, Max: {test_df['utterance_length'].max():.2f}")
```

=== Utterance Length Statistics ===
Training set - Mean: 7.95, Median: 6, Max: 69
Dev set - Mean: 7.91, Median: 6, Max: 37
Test set - Mean: 8.21, Median: 7, Max: 45


9. Speaker Analysis

```
In [ ]: # Analyze speaker distribution
all_speakers = pd.concat([train_df['Speaker'], dev_df['Speaker'], test_df['Speaker']])
speaker_counts = all_speakers.value_counts()

plt.figure(figsize=(10, 6))
speaker_counts.plot(kind='bar')
plt.title('Distribution of Utterances by Speaker', fontsize=14)
plt.xlabel('Speaker')
plt.ylabel('Number of Utterances')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

print(f"Total unique speakers: {len(speaker_counts)}")
print("\n=== Speaker Distribution ===")
print(speaker_counts)
```


Distribution of Utterances by Speaker



Total unique speakers: 304

=== Speaker Distribution ===

Speaker

Joey 2070

Ross 2048

Rachel 1955

Phoebe 1797

Monica 1782

...

Waiter 1

Cookie 1

Passerby 1

Petrie 1

Sergei 1

Name: count, Length: 304, dtype: int64

10. Emotion-Sentiment Correlation

```
In [ ]: # Analyze correlation between emotion and sentiment
def create_emotion_sentiment_heatmap(df, title):
    emotion_sentiment_crosstab = pd.crosstab(df['Emotion'], df['Sentiment'])

    plt.figure(figsize=(10, 8))
    sns.heatmap(emotion_sentiment_crosstab, annot=True, fmt='d', cmap='Blues')
    plt.title(f'Emotion vs Sentiment Correlation - {title}', fontsize=14)
    plt.xlabel('Sentiment')
```


```
plt.ylabel('Emotion')
plt.tight_layout()
plt.show()

return emotion_sentiment_crosstab

print("=== Training Set Emotion-Sentiment Correlation ===")
train_crosstab = create_emotion_sentiment_heatmap(train_df, 'Training Set')

=== Training Set Emotion-Sentiment Correlation ===
```

Emotion vs Sentiment Correlation - Training Set





11. Dialogue Context Analysis

```
In [ ]: # Analyze dialogue structure
dialogue_counts = train_df['Dialogue_ID'].value_counts()
print(f"Total dialogues in training set: {len(dialogue_counts)}")
print(f"Average utterances per dialogue: {dialogue_counts.mean():.2f}")
print(f"Min utterances in a dialogue: {dialogue_counts.min()}")
print(f"Max utterances in a dialogue: {dialogue_counts.max()}")

# Plot dialogue length distribution
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.hist(dialogue_counts.values, bins=30)
plt.xlabel('Number of Utterances per Dialogue')
plt.ylabel('Frequency')
plt.title('Distribution of Dialogue Lengths')

plt.subplot(1, 2, 2)
plt.boxplot(dialogue_counts.values)
plt.ylabel('Number of Utterances')
plt.title('Dialogue Length Box Plot')
plt.tight_layout()
plt.show()
```

Total dialogues in training set: 1038
Average utterances per dialogue: 9.62
Min utterances in a dialogue: 1
Max utterances in a dialogue: 24



12. Sample Conversations Analysis

```
In [ ]: # Display sample conversations with emotions
def display_sample_dialogue(df, dialogue_id, max_utterances=10):
    dialogue = df[df['Dialogue_ID'] == dialogue_id].head(max_utterances)
    print(f"\n=== Sample Dialogue (ID: {dialogue_id}) ===")
    for idx, row in dialogue.iterrows():
```

```
print(f"{row['Speaker']}: {row['Utterance']}")  
print(f"    Emotion: {row['Emotion']}, Sentiment: {row['Sentiment']}")  
print()
```

Display a few sample dialogues

```
sample_dialogue_ids = train_df['Dialogue_ID'].unique()[:3]  
for dialogue_id in sample_dialogue_ids:  
    display_sample_dialogue(train_df, dialogue_id)
```

=== Sample Dialogue (ID: 0) ===

Chandler: also I was the point person on my company's transition from the KL-5 to GR-6 system.

Emotion: neutral, Sentiment: neutral

The Interviewer: You must've had your hands full.

Emotion: neutral, Sentiment: neutral

Chandler: That I did. That I did.

Emotion: neutral, Sentiment: neutral

The Interviewer: So let's talk a little bit about your duties.

Emotion: neutral, Sentiment: neutral

Chandler: My duties? All right.

Emotion: surprise, Sentiment: positive

The Interviewer: Now you'll be heading a whole division, so you'll have a lot of duties.

Emotion: neutral, Sentiment: neutral

Chandler: I see.

Emotion: neutral, Sentiment: neutral

The Interviewer: But there'll be perhaps 30 people under you so you can dump a certain amount on them.

Emotion: neutral, Sentiment: neutral

Chandler: Good to know.

Emotion: neutral, Sentiment: neutral

The Interviewer: We can go into detail

Emotion: neutral, Sentiment: neutral

=== Sample Dialogue (ID: 1) ===

Joey: But then who? The waitress I went out with last month?

Emotion: surprise, Sentiment: negative

Rachel: You know? Forget it!

Emotion: sadness, Sentiment: negative

Joey: No-no-no-no, no! Who, who were you talking about?

Emotion: surprise, Sentiment: negative

Rachel: No, I-I-I-I don't, I actually don't know

Emotion: fear, Sentiment: negative

Joey: Ok!

Emotion: neutral, Sentiment: neutral

Joey: All right, well...

Emotion: neutral, Sentiment: neutral

Rachel: Yeah, sure!

Emotion: neutral, Sentiment: neutral

=== Sample Dialogue (ID: 2) ===

Chandler: Hey, Mon.

Emotion: neutral, Sentiment: neutral

Monica: Hey-hey-hey. You wanna hear something that sucks.

Emotion: neutral, Sentiment: neutral

Chandler: Do I ever.

Emotion: joy, Sentiment: positive

Monica: Chris says they're closing down the bar.

Emotion: sadness, Sentiment: negative

Chandler: No way!

Emotion: surprise, Sentiment: negative

Monica: Yeah, apparently they're turning it into some kinda coffee place.

Emotion: neutral, Sentiment: neutral

Chandler: Just coffee! Where are we gonna hang out now?

Emotion: disgust, Sentiment: negative

Monica: Got me.

Emotion: sadness, Sentiment: negative

Chandler: Can I get a beer.

Emotion: neutral, Sentiment: neutral

Monica: Hey, did you pick a roommate?

Emotion: neutral, Sentiment: neutral

13. Word Frequency Analysis

```
In [ ]: from collections import Counter
import re

# Function to clean and tokenize text
def tokenize(text):
    # Convert to lowercase and remove punctuation
    text = str(text).lower()
    text = re.sub(r'^\w\s', '', text)
    return text.split()

# Get word frequencies for each emotion
emotion_words = {}
for emotion in train_df['Emotion'].unique():
    emotion utterances = train_df[train_df['Emotion'] == emotion]['Utterance']
    all_words = []
    for utterance in emotion utterances:
        all_words.extend(tokenize(utterance))
    emotion_words[emotion] = Counter(all_words).most_common(20)

In [ ]: # Display top words for each emotion
print("=== Top 10 Words per Emotion ===")
for emotion, word_counts in emotion_words.items():
```

```
print(f"\n{emotion}:")  
for word, count in word_counts[:10]:  
    print(f"    {word}: {count}")
```

=== Top 10 Words per Emotion ===

neutral:

you: 1373
i: 1318
the: 916
to: 783
a: 773
and: 538
that: 457
it: 449
okay: 425
yeah: 406

surprise:

you: 387
what: 254
oh: 238
i: 211
the: 150
my: 138
a: 124
that: 122
god: 114
are: 105

fear:

i: 132
you: 82

to: 57
a: 45
no: 44
dont: 39
oh: 37
im: 37
and: 33
the: 33

sadness:

i: 364
you: 190
to: 162
the: 143
a: 126
im: 120
and: 119
sorry: 97
it: 95
me: 92

joy:

i: 535
you: 499
the: 287
to: 251
a: 249
oh: 232
and: 223

it: 216
so: 164
hey: 143

disgust:
you: 96
i: 78
the: 74
a: 69
to: 51
and: 44
that: 44
oh: 41
no: 35
it: 33

anger:
you: 431
i: 374
the: 247
to: 218
and: 174
a: 168
it: 157
that: 143
me: 137
no: 130

14. Emotion Transition Analysis

```
In [ ]: # Analyze emotion transitions within dialogues
def analyze_emotion_transitions(df):
    transitions = []

    for dialogue_id in df['Dialogue_ID'].unique():
        dialogue = df[df['Dialogue_ID'] == dialogue_id].sort_values('Utterance_ID')
        emotions = dialogue['Emotion'].tolist()

        for i in range(len(emotions) - 1):
            transitions.append((emotions[i], emotions[i+1]))

    return Counter(transitions)
```

```
In [ ]: # Get emotion transitions
transitions = analyze_emotion_transitions(train_df)
top_transitions = transitions.most_common(15)

print("=== Top 15 Emotion Transitions ===")
for (from_emotion, to_emotion), count in top_transitions:
    print(f"{from_emotion} → {to_emotion}: {count}")
```


=== Top 15 Emotion Transitions ===

neutral → neutral: 2354

neutral → joy: 631

joy → neutral: 602

joy → joy: 520

neutral → surprise: 500

surprise → neutral: 492

anger → neutral: 358

neutral → anger: 352

anger → anger: 314

neutral → sadness: 212

sadness → neutral: 206

surprise → surprise: 179

surprise → joy: 167

joy → surprise: 165

sadness → sadness: 147

```
In [ ]: # Create transition matrix
unique_emotions = sorted(train_df['Emotion'].unique())
transition_matrix = pd.DataFrame(0, index=unique_emotions, columns=unique_emotions)

for (from_emotion, to_emotion), count in transitions.items():
    transition_matrix.loc[from_emotion, to_emotion] = count
```

```
In [ ]: # Normalize by row to get probabilities
transition_prob = transition_matrix.div(transition_matrix.sum(axis=1), axis=0)

plt.figure(figsize=(10, 8))
sns.heatmap(transition_prob, annot=True, fmt='.2f', cmap='YlOrRd')
```

```
plt.title('Emotion Transition Probability Matrix')  
plt.xlabel('To Emotion')  
plt.ylabel('From Emotion')  
plt.tight_layout()  
plt.show()
```





15. Summary Statistics and Insights

```
In [ ]: print("=== MELD Dataset Summary ===")
print(f"Total utterances: {total_utterances}")
print(f"Total unique dialogues: {len(pd.concat([train_df['Dialogue_ID'], dev_df['Dialogue_ID']]))}")
print(f"Average utterance length: {pd.concat([train_df['utterance_length'], dev_df['utterance_length']]).mean():.2f}")
print(f"\nEmotion classes: {sorted(train_df['Emotion'].unique())}")
print(f"Sentiment classes: {sorted(train_df['Sentiment'].unique())}")

# Class imbalance analysis
print("\n=== Class Imbalance Analysis ===")
emotion_imbalance = emotion_dist.max() / emotion_dist.min()
sentiment_imbalance = sentiment_dist.max() / sentiment_dist.min()
print(f"Emotion class imbalance ratio: {emotion_imbalance:.2f}")
print(f"Sentiment class imbalance ratio: {sentiment_imbalance:.2f}")
```

=== MELD Dataset Summary ===

Total utterances: 13708

Total unique dialogues: 1039

Average utterance length: 8.00 words

Emotion classes: ['anger', 'disgust', 'fear', 'joy', 'neutral', 'sadness', 'surprise']

Sentiment classes: ['negative', 'neutral', 'positive']

=== Class Imbalance Analysis ===

Emotion class imbalance ratio: 17.98

Sentiment class imbalance ratio: 2.08

```
In [ ]: # Key insights for model development
print("\n=== Key Insights for Model Development ===")
print("1. The dataset shows significant class imbalance, especially in emotions")
print("2. Neutral emotion is dominant, which might affect model performance")
print("3. Average utterance length is relatively short, suitable for transformer models")
print("4. Strong correlation between certain emotions and sentiments")
print("5. Context from dialogue flow could be important for emotion recognition")
```

=== Key Insights for Model Development ===

1. The dataset shows significant class imbalance, especially in emotions

2. Neutral emotion is dominant, which might affect model performance

3. Average utterance length is relatively short, suitable for transformer models

4. Strong correlation between certain emotions and sentiments

5. Context from dialogue flow could be important for emotion recognition

16. Prepare Data for Transformer Models

```
In [ ]: # Create a function to prepare data for transformer models
def prepare_for_transformers(df):
    """
    Prepare the dataset for transformer-based models
    """

    # Create emotion to index mapping
    emotion_to_idx = {emotion: idx for idx, emotion in enumerate(sorted(df['Emotion'].unique()))}
    sentiment_to_idx = {sentiment: idx for idx, sentiment in enumerate(sorted(df['Sentiment'].unique()))}

    # Add numerical labels
    df['emotion_label'] = df['Emotion'].map(emotion_to_idx)
    df['sentiment_label'] = df['Sentiment'].map(sentiment_to_idx)

    # Group by dialogue for context modeling
    dialogues = []
    for dialogue_id in df['Dialogue_ID'].unique():
        dialogue = df[df['Dialogue_ID'] == dialogue_id].sort_values('Utterance_ID')
        dialogues.append({
            'dialogue_id': dialogue_id,
            'utterances': dialogue['Utterance'].tolist(),
            'speakers': dialogue['Speaker'].tolist(),
            'emotions': dialogue['emotion_label'].tolist(),
            'sentiments': dialogue['sentiment_label'].tolist()
        })
```

```
return dialogues, emotion_to_idx, sentiment_to_idx
```

```
In [ ]: # Prepare training data
train_dialogues, emotion_to_idx, sentiment_to_idx = prepare_for_transformers(1
print(f"Prepared {len(train_dialogues)} dialogues for training")
print(f"\nEmotion mapping: {emotion_to_idx}")
print(f"\nSentiment mapping: {sentiment_to_idx}")
```

Prepared 1038 dialogues for training

Emotion mapping: {'anger': 0, 'disgust': 1, 'fear': 2, 'joy': 3, 'neutral': 4, 'sadness': 5, 'surprise': 6}

Sentiment mapping: {'negative': 0, 'neutral': 1, 'positive': 2}

```
In [ ]: # Save mappings for later use
import json
if not os.path.exists('data'):
    os.makedirs('data')

mappings = {
    'emotion_to_idx': emotion_to_idx,
    'sentiment_to_idx': sentiment_to_idx,
    'idx_to_emotion': {v: k for k, v in emotion_to_idx.items()},
    'idx_to_sentiment': {v: k for k, v in sentiment_to_idx.items()}
}

with open('data/label_mappings.json', 'w') as f:
```

```
    json.dump(mappings, f, indent=2)
print("\nLabel mappings saved to 'data/label_mappings.json'")
```

Label mappings saved to 'data/label_mappings.json'

Conclusion

This EDA reveals several important insights for building emotion recognition models:

1. **Class Imbalance:** The dataset shows significant imbalance, particularly with "neutral" emotion being dominant. Consider using weighted loss functions or resampling techniques.
2. **Multi-modal Nature:** While we focused on text, MELD includes audio and visual features that could enhance model performance.
3. **Context Importance:** Emotions often depend on dialogue context, making this dataset ideal for context-aware transformer models.
4. **Short Utterances:** Most utterances are relatively short (median ~10 words), which is suitable for transformer architectures.
5. **Speaker Patterns:** Different speakers show distinct emotion patterns, which could be leveraged in multi-speaker models.

Next steps:

- Implement transformer-based models (BERT, RoBERTa, etc.) for emotion classification
- Experiment with context-aware architectures that consider dialogue history
- Apply techniques to handle class imbalance
- Consider multi-task learning for joint emotion and sentiment prediction

Model Implementation

Implementation of LSTM, BERT, RoBERTa with context-aware models using TensorFlow

1. Import Required Libraries and Setup

```
In [ ]: import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, Model
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLrOnPlateau
import tensorflow_hub as hub
import tensorflow_text as text
from transformers import TFAutoModel, AutoTokenizer, TFBertModel, TFRobertaModel
from sklearn.metrics import classification_report, confusion_matrix, f1_score
```

```
from sklearn.utils.class_weight import compute_class_weight
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
import json
import pickle
import warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: # Set random seeds
def set_seed(seed=42):
    np.random.seed(seed)
    tf.random.set_seed(seed)

set_seed(42)
```

```
In [ ]: # Check GPU availability
print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
if len(tf.config.list_physical_devices('GPU')) > 0:
    print("GPU is available")
    # Enable mixed precision for better performance
    from tensorflow.keras.mixed_precision import set_global_policy
    set_global_policy('mixed_float16')
```

Num GPUs Available: 1
GPU is available

2. Data Preprocessing and Dataset Classes

```
In [ ]: import numpy as np
import tensorflow as tf
from transformers import AutoTokenizer

class MELDDataGeneratorBERT(tf.keras.utils.Sequence):
    """Fixed data generator for MELD dataset with proper tokenization"""
    def __init__(self, dialogues, tokenizer, batch_size=32, max_length=128,
                 context_window=3, shuffle=True):
        self.dialogues = dialogues
        self.tokenizer = tokenizer
        self.batch_size = batch_size
        self.max_length = max_length
        self.context_window = context_window
        self.shuffle = shuffle

        # IMPORTANT: Set tokenizer max length
        self.tokenizer.model_max_length = max_length

        # Prepare data
        self.utterances = []
        self.contexts = []
        self.emotions = []
        self.sentiments = []
        self._prepare_data()
```

```

self.indices = np.arange(len(self.utterances))
if self.shuffle:
    np.random.shuffle(self.indices)

def _prepare_data(self):
    """Flatten dialogues into utterances with context"""
    for dialogue in self.dialogues:
        utterances = dialogue['utterances']
        emotions = dialogue['emotions']
        sentiments = dialogue['sentiments']
        speakers = dialogue['speakers']

        for i in range(len(utterances)):
            # Current utterance - ensure it's a string
            self.utterances.append(str(utterances[i]))
            # Ensure emotions and sentiments are integers
            self.emotions.append(int(emotions[i]))
            self.sentiments.append(int(sentiments[i]))

            # Context (previous utterances)
            context = []
            for j in range(max(0, i - self.context_window), i):
                context.append(f"{speakers[j]}: {utterances[j]}")
            self.contexts.append(" [SEP] ".join(context) if context else '

def __len__(self):
    return int(np.ceil(len(self.utterances) / self.batch_size))

```

```

def __getitem__(self, idx):
    # Get batch indices
    start_idx = idx * self.batch_size
    end_idx = min((idx + 1) * self.batch_size, len(self.utterances))
    batch_indices = self.indices[start_idx:end_idx]

    batch_texts = []
    batch_emotions = []
    batch_sentiments = []

    for i in batch_indices:
        # Combine context and utterance
        if self.contexts[i]:
            text = f"{self.contexts[i]} [SEP] {self.utterances[i]}"
        else:
            text = self.utterances[i]

        batch_texts.append(text)
        batch_emotions.append(self.emotions[i])
        batch_sentiments.append(self.sentiments[i])

    # Tokenize batch with explicit parameters
    encoded = self.tokenizer(
        batch_texts,
        padding='max_length',
        truncation=True,
        max_length=self.max_length,
        return_tensors='np',
        return_attention_mask=True,

```

```

        return_token_type_ids=False # BERT might return this, we don't need it
    )

    # Ensure we have the right shape
    input_ids = encoded['input_ids']
    attention_mask = encoded['attention_mask']

    # Pad manually if needed (shouldn't be necessary but just in case)
    if input_ids.shape[1] < self.max_length:
        pad_length = self.max_length - input_ids.shape[1]
        input_ids = np.pad(input_ids, ((0, 0), (0, pad_length)),
                           constant_values=self.tokenizer.pad_token_id)
        attention_mask = np.pad(attention_mask, ((0, 0), (0, pad_length)),
                                constant_values=0)
    elif input_ids.shape[1] > self.max_length:
        input_ids = input_ids[:, :self.max_length]
        attention_mask = attention_mask[:, :self.max_length]

    # Convert to proper numpy arrays with correct dtypes
    input_ids = np.array(input_ids, dtype=np.int32)
    attention_mask = np.array(attention_mask, dtype=np.int32)

    # Return dictionaries with numpy arrays
    return {
        'input_ids': input_ids,
        'attention_mask': attention_mask
    }, {
        'emotion': np.array(batch_emotions, dtype=np.int32),
        'sentiment': np.array(batch_sentiments, dtype=np.int32)
    }

```

```

    }

    def on_epoch_end(self):
        """Shuffle indices after each epoch"""
        if self.shuffle:
            np.random.shuffle(self.indices)

```

```

In [ ]: class MELDDataGenerator(tf.keras.utils.Sequence):
        """Fixed data generator with proper max_length padding"""

        def __init__(self, dialogues, tokenizer, batch_size=32, max_length=128,
                      context_window=3, shuffle=True):
            self.dialogues = dialogues
            self.tokenizer = tokenizer
            self.batch_size = batch_size
            self.max_length = max_length
            self.context_window = context_window
            self.shuffle = shuffle

            # Prepare data
            self.utterances = []
            self.contexts = []
            self.emotions = []
            self.sentiments = []
            self._prepare_data()

            self.indices = np.arange(len(self.utterances))
            if self.shuffle:
                np.random.shuffle(self.indices)

```

```

def _prepare_data(self):
    """Flatten dialogues into utterances with context"""
    for dialogue in self.dialogues:
        utterances = dialogue['utterances']
        emotions = dialogue['emotions']
        sentiments = dialogue['sentiments']
        speakers = dialogue['speakers']

        for i in range(len(utterances)):
            # Current utterance
            self.utterances.append(utterances[i])
            self.emotions.append(emotions[i])
            self.sentiments.append(sentiments[i])

            # Context (previous utterances)
            context = []
            for j in range(max(0, i - self.context_window), i):
                context.append(f"{speakers[j]}: {utterances[j]}")
            self.contexts.append(" [SEP] ".join(context) if context else ' ')

def __len__(self):
    return int(np.ceil(len(self.utterances) / self.batch_size))

def __getitem__(self, idx):
    batch_indices = self.indices[idx * self.batch_size:(idx + 1) * self.batch_size]

    batch_texts = []
    batch_emotions = []

```



```

batch_sentiments = []

for i in batch_indices:
    # Combine context and utterance
    if self.contexts[i]:
        text = f"{self.contexts[i]} [SEP] {self.utterances[i]}"
    else:
        text = self.utterances[i]

    batch_texts.append(text)
    batch_emotions.append(self.emotions[i])
    batch_sentiments.append(self.sentiments[i])

# Tokenize batch with FIXED padding to max_length
encoded = self.tokenizer(
    batch_texts,
    padding='max_length', # Changed from padding=True
    truncation=True,
    max_length=self.max_length,
    return_tensors='tf'
)

return {
    'input_ids': encoded['input_ids'],
    'attention_mask': encoded['attention_mask']
}, {
    'emotion': tf.convert_to_tensor(batch_emotions),
    'sentiment': tf.convert_to_tensor(batch_sentiments)
}

```

```
def on_epoch_end(self):
    if self.shuffle:
        np.random.shuffle(self.indices)
```

```
In [ ]: # Test function to verify the generator works correctly
def test_data_generator(dialogues, tokenizer, batch_size=16, max_length=128):
    """Test the data generator to ensure it works correctly"""

    print("Testing data generator...")

    # Create generator
    generator = MELDDDataGenerator(
        dialogues,
        tokenizer,
        batch_size=batch_size,
        max_length=max_length,
        shuffle=False
    )

    print(f"Total batches: {len(generator)}")
    print(f"Total samples: {len(generator.utterances)}")

    # Test first batch
    batch_x, batch_y = generator[0]

    print(f"\nFirst batch shapes:")
    print(f"  input_ids: {batch_x['input_ids'].shape}")
    print(f"  attention_mask: {batch_x['attention_mask'].shape}")
```

```

print(f"  emotions: {batch_y['emotion'].shape}")
print(f"  sentiments: {batch_y['sentiment'].shape}")

# Verify shapes
actual_batch_size = batch_x['input_ids'].shape[0]
assert actual_batch_size <= batch_size, f"Batch size too large! Expected <
assert batch_x['input_ids'].shape[1] == max_length, f"Max length mismatch
assert batch_x['attention_mask'].shape == batch_x['input_ids'].shape, "Att
assert batch_y['emotion'].shape[0] == actual_batch_size, "Emotion batch si
assert batch_y['sentiment'].shape[0] == actual_batch_size, "Sentiment batch

print("\n✓ Data generator test passed!")

# Print sample data
print("\nSample from first batch:")
print(f"  Text (decoded): {tokenizer.decode(batch_x['input_ids'][0], skip_
print(f"  Emotion label: {batch_y['emotion'][0]}")
print(f"  Sentiment label: {batch_y['sentiment'][0]}")

return generator

```

```

In [ ]: def create_tf_dataset(dialogues, tokenizer, batch_size=32, max_length=128,
                                context_window=3, shuffle=True):
    """Create TensorFlow dataset from dialogues"""
    utterances = []
    contexts = []
    emotions = []
    sentiments = []

```

```

# Prepare data
for dialogue in dialogues:
    utts = dialogue['utterances']
    emos = dialogue['emotions']
    sents = dialogue['sentiments']
    speakers = dialogue['speakers']

    for i in range(len(utts)):
        utterances.append(utts[i])
        emotions.append(emos[i])
        sentiments.append(sents[i])

    # Context
    context = []
    for j in range(max(0, i - context_window), i):
        context.append(f"{speakers[j]}: {utts[j]}")
    contexts.append(" [SEP] ".join(context) if context else "")

# Create dataset from tensors
dataset = tf.data.Dataset.from_tensor_slices((utterances, contexts, emotions))

if shuffle:
    dataset = dataset.shuffle(buffer_size=1000)

# Map function to combine context and utterance and tokenize
def map_fn(utterance, context, emotion, sentiment):
    text = tf.cond(tf.equal(tf.strings.length(context), 0),
                    lambda: utterance,
                    lambda: tf.strings.join([context, "[SEP]", utterance],

```

```

# Tokenize
encoded = tokenizer(
    text,
    padding='max_length',
    truncation=True,
    max_length=max_length,
    return_tensors='tf'
)

return {
    'input_ids': tf.squeeze(encoded['input_ids'], axis=0),
    'attention_mask': tf.squeeze(encoded['attention_mask'], axis=0)
}, {
    'emotion': tf.cast(emotion, tf.int32),
    'sentiment': tf.cast(sentiment, tf.int32)
}

dataset = dataset.map(map_fn, num_parallel_calls=tf.data.AUTOTUNE)

# Batch the dataset
dataset = dataset.batch(batch_size)

# Set the shapes of the batched tensors
dataset = dataset.map(
    lambda x, y: (
        {'input_ids': tf.ensure_shape(x['input_ids'], (None, max_length)),
         'attention_mask': tf.ensure_shape(x['attention_mask'], (None, max_length))},
        {'emotion': tf.ensure_shape(y['emotion'], (None,)),
         'sentiment': tf.ensure_shape(y['sentiment'], (None,))}
    )
)

```

```

        'sentiment': tf.ensure_shape(y['sentiment'], (None,))}
    ),
    num_parallel_calls=tf.data.AUTOTUNE
)

dataset = dataset.prefetch(tf.data.AUTOTUNE)

return dataset

```

```

In [ ]: def create_tf_dataset_2(dialogues, tokenizer, batch_size=32, max_length=128,
                                context_window=3, shuffle=True):
    """Create TensorFlow dataset from dialogues"""
    utterances = []
    contexts = []
    emotions = []
    sentiments = []

    # Prepare data
    for dialogue in dialogues:
        utts = dialogue['utterances']
        emos = dialogue['emotions']
        sents = dialogue['sentiments']
        speakers = dialogue['speakers']

        for i in range(len(utts)):
            utterances.append(utts[i])
            emotions.append(emos[i])
            sentiments.append(sents[i])

```

```

    # Context
    context = []
    for j in range(max(0, i - context_window), i):
        context.append(f"{speakers[j]}: {utts[j]}")
    contexts.append(" [SEP] ".join(context) if context else "")

# Create dataset
def generator():
    indices = np.arange(len(utterances))
    if shuffle:
        np.random.shuffle(indices)

    for i in indices:
        text = f"{contexts[i]} [SEP] {utterances[i]}" if contexts[i] else
        yield text, emotions[i], sentiments[i]

# Define output signature
output_signature = (
    tf.TensorSpec(shape=(), dtype=tf.string),
    tf.TensorSpec(shape=(), dtype=tf.int32),
    tf.TensorSpec(shape=(), dtype=tf.int32)
)

dataset = tf.data.Dataset.from_generator(
    generator,
    output_signature=output_signature
)

```

```

# Tokenize and batch
def tokenize_batch(texts, emotions, sentiments):
    # This will be called with batched data
    encoded = tokenizer(
        texts.numpy().tolist(),
        padding=True,
        truncation=True,
        max_length=max_length,
        return_tensors='tf'
    )

    return {
        'input_ids': encoded['input_ids'],
        'attention_mask': encoded['attention_mask'],
        'emotion_labels': emotions,
        'sentiment_labels': sentiments
    }

dataset = dataset.batch(batch_size)
dataset = dataset.map(
    lambda x, y, z: tf.py_function(
        tokenize_batch,
        [x, y, z],
        {
            'input_ids': tf.int32,
            'attention_mask': tf.int32,
            'emotion_labels': tf.int32,
            'sentiment_labels': tf.int32
        }
    )
)

```



```

    ),
    num_parallel_calls=tf.data.AUTOTUNE
)

if shuffle:
    dataset = dataset.shuffle(buffer_size=1000)

dataset = dataset.prefetch(tf.data.AUTOTUNE)

return dataset

```

3. Model Architectures

3.1 LSTM-based Model

```

In [ ]: class LSTMEmotionModel(tf.keras.Model):
        """LSTM model for emotion and sentiment classification"""
        def __init__(self, vocab_size, embedding_dim=300, lstm_units=256,
                        num_emotions=7, num_sentiments=3, dropout_rate=0.3):
            super(LSTMEmotionModel, self).__init__()

            # Embedding layer
            self.embedding = layers.Embedding(vocab_size, embedding_dim)

            # LSTM layers
            self.lstm1 = layers.Bidirectional(
                layers.LSTM(lstm_units, return_sequences=True, dropout=dropout_rate)

```

```

    )
    self.lstm2 = layers.Bidirectional(
        layers.LSTM(lstm_units, dropout=dropout_rate)
    )

    # Dropout
    self.dropout = layers.Dropout(dropout_rate)

    # Task-specific heads
    self.emotion_dense = layers.Dense(128, activation='relu')
    self.emotion_output = layers.Dense(num_emotions, name='emotion')

    self.sentiment_dense = layers.Dense(128, activation='relu')
    self.sentiment_output = layers.Dense(num_sentiments, name='sentiment')

def call(self, inputs, training=False):
    # Embedding
    x = self.embedding(inputs['input_ids'])

    # LSTM encoding
    x = self.lstm1(x, training=training)
    x = self.lstm2(x, training=training)

    # Dropout
    x = self.dropout(x, training=training)

    # Task-specific predictions
    emotion_features = self.emotion_dense(x)
    emotion_logits = self.emotion_output(emotion_features)

```

```

sentiment_features = self.sentiment_dense(x)
sentiment_logits = self.sentiment_output(sentiment_features)

return {
    'emotion': emotion_logits,
    'sentiment': sentiment_logits
}

```

3.2 BERT-based Multi-Task Model

```

In [ ]: # Disable mixed precision for inference to avoid dtype issues
tf.keras.mixed_precision.set_global_policy('float32')

class BERTMultiTaskModel(tf.keras.Model):
    """BERT model for multi-task emotion and sentiment classification"""
    def __init__(self, model_name='bert-base-uncased', num_emotions=7,
                  num_sentiments=3, dropout_rate=0.3):
        super(BERTMultiTaskModel, self).__init__()

        # BERT encoder
        self.bert = TFAutoModel.from_pretrained(model_name)
        hidden_size = self.bert.config.hidden_size

        # Freeze BERT layers initially (optional)
        self.bert.trainable = True

        # Dropout

```

```

self.dropout = tf.keras.layers.Dropout(dropout_rate)

# Emotion classification head
self.emotion_dense1 = tf.keras.layers.Dense(hidden_size, activation='tanh')
self.emotion_dropout = tf.keras.layers.Dropout(dropout_rate)
self.emotion_dense2 = tf.keras.layers.Dense(hidden_size // 2, activation='tanh')
self.emotion_output = tf.keras.layers.Dense(num_emotions, name='emotion_output')

# Sentiment classification head
self.sentiment_dense1 = tf.keras.layers.Dense(hidden_size, activation='tanh')
self.sentiment_dropout = tf.keras.layers.Dropout(dropout_rate)
self.sentiment_dense2 = tf.keras.layers.Dense(hidden_size // 2, activation='tanh')
self.sentiment_output = tf.keras.layers.Dense(num_sentiments, name='sentiment_output')

def call(self, inputs, training=False):
    # Ensure input_ids are int32
    input_ids = tf.cast(inputs['input_ids'], tf.int32)
    attention_mask = tf.cast(inputs['attention_mask'], tf.int32)

    # BERT encoding
    bert_outputs = self.bert(
        input_ids=input_ids,
        attention_mask=attention_mask,
        training=training
    )

    # Use pooled output (CLS token)
    pooled_output = bert_outputs.pooler_output
    pooled_output = self.dropout(pooled_output, training=training)

```

```

# Emotion prediction
emotion_x = self.emotion_dense1(pooled_output)
emotion_x = self.emotion_dropout(emotion_x, training=training)
emotion_x = self.emotion_dense2(emotion_x)
emotion_logits = self.emotion_output(emotion_x)

# Sentiment prediction
sentiment_x = self.sentiment_dense1(pooled_output)
sentiment_x = self.sentiment_dropout(sentiment_x, training=training)
sentiment_x = self.sentiment_dense2(sentiment_x)
sentiment_logits = self.sentiment_output(sentiment_x)

return {
    'emotion': emotion_logits,
    'sentiment': sentiment_logits
}

```

3.3 Context-Aware DialogueRNN Model

```

In [ ]: class DialogueRNN(tf.keras.Model):
        """DialogueRNN for context-aware emotion recognition"""
        def __init__(self, base_model_name='bert-base-uncased', hidden_dim=128,
                        num_emotions=7, num_sentiments=3, dropout_rate=0.3):
            super(DialogueRNN, self).__init__()

        # Base encoder (BERT or RoBERTa)
        self.encoder = TFAutoModel.from_pretrained(base_model_name)

```

```

encoder_hidden_size = self.encoder.config.hidden_size

# Global GRU for context
self.global_gru = layers.Bidirectional(
    layers.GRU(hidden_dim, return_sequences=True)
)

# Emotion GRU
self.emotion_gru = layers.GRU(hidden_dim, return_sequences=False)

# Attention mechanism
self.attention = layers.MultiHeadAttention(
    num_heads=8, key_dim=hidden_dim
)

# Fusion layer
self.fusion = layers.Dense(hidden_dim, activation='relu')

# Dropout
self.dropout = layers.Dropout(dropout_rate)

# Classification heads
self.emotion_output = layers.Dense(num_emotions, name='emotion')
self.sentiment_output = layers.Dense(num_sentiments, name='sentiment')

def call(self, inputs, training=False):
    # Ensure proper dtypes
    input_ids = tf.cast(inputs['input_ids'], tf.int32)
    attention_mask = tf.cast(inputs['attention_mask'], tf.int32)

```

```
# Encode utterance
encoded = self.encoder(
    input_ids=input_ids,
    attention_mask=attention_mask,
    training=training
)

# Get CLS token representation
utterance_features = encoded.pooler_output

# Add sequence dimension for GRU
utterance_features = tf.expand_dims(utterance_features, axis=1)

# Global context (simplified for single utterance inference)
global_context = self.global_gru(utterance_features, training=training)

# Apply attention
attended_features = self.attention(
    utterance_features, global_context,
    training=training
)

# Emotion modeling
emotion_features = self.emotion_gru(attended_features, training=training)
emotion_features = self.dropout(emotion_features, training=training)

# Fusion
final_features = self.fusion(emotion_features)
```

```

# Predictions
emotion_logits = self.emotion_output(final_features)
sentiment_logits = self.sentiment_output(final_features)

return {
    'emotion': emotion_logits,
    'sentiment': sentiment_logits
}

```

3.4 COSMIC-style Model with Attention

```

In [ ]: class COSMICModel(tf.keras.Model):
        """COSMIC-style model with commonsense reasoning"""
        def __init__(self, model_name='roberta-base', num_emotions=7,
                      num_sentiments=3, dropout_rate=0.3):
            super(COSMICModel, self).__init__()

            # Main encoder
            self.encoder = TFAutoModel.from_pretrained(model_name)
            hidden_size = self.encoder.config.hidden_size

            # Commonsense encoder (can be same or different model)
            self.commonsense_encoder = TFAutoModel.from_pretrained('bert-base-uncased')

            # Cross-attention for fusion
            self.cross_attention = layers.MultiHeadAttention(
                num_heads=8, key_dim=hidden_size // 8
            )

```



```

)

# Fusion layers
self.fusion_dense1 = layers.Dense(hidden_size, activation='relu')
self.fusion_dropout = layers.Dropout(dropout_rate)
self.fusion_dense2 = layers.Dense(hidden_size // 2, activation='relu')

# Self-attention
self.self_attention = layers.MultiHeadAttention(
    num_heads=8, key_dim=hidden_size // 16
)

# Classification heads
self.emotion_output = layers.Dense(num_emotions, name='emotion')
self.sentiment_output = layers.Dense(num_sentiments, name='sentiment')

self.dropout = layers.Dropout(dropout_rate)

def call(self, inputs, commonsense_inputs=None, training=False):
    input_ids = tf.cast(inputs['input_ids'], tf.int32)
    attention_mask = tf.cast(inputs['attention_mask'], tf.int32)
    # Encode utterance
    utterance_outputs = self.encoder(
        input_ids=input_ids,
        attention_mask=attention_mask,
        training=training
    )
    utterance_features = utterance_outputs.pooler_output

```

```

if commonsense_inputs is not None:
    # Encode commonsense
    commonsense_outputs = self.commonsense_encoder(
        commonsense_inputs['input_ids'],
        attention_mask=commonsense_inputs['attention_mask'],
        training=training
    )
    commonsense_features = commonsense_outputs.pooler_output

    # Cross-attention fusion
    utterance_features = tf.expand_dims(utterance_features, axis=1)
    commonsense_features = tf.expand_dims(commonsense_features, axis=1)

    fused_features = self.cross_attention(
        utterance_features,
        commonsense_features,
        training=training
    )
    fused_features = tf.squeeze(fused_features, axis=1)

    # Further fusion
    fused_features = self.fusion_dense1(fused_features)
    fused_features = self.fusion_dropout(fused_features, training=training)
    final_features = self.fusion_dense2(fused_features)
else:
    final_features = utterance_features

final_features = self.dropout(final_features, training=training)

```

```

# Predictions
emotion_logits = self.emotion_output(final_features)
sentiment_logits = self.sentiment_output(final_features)

return {
    'emotion': emotion_logits,
    'sentiment': sentiment_logits
}

```

4. Loss Functions and Metrics

```

In [ ]: class MultiTaskLoss(tf.keras.losses.Loss):
        """Multi-task loss with task weighting"""
        def __init__(self, emotion_weight=1.0, sentiment_weight=0.5,
                      emotion_class_weights=None, sentiment_class_weights=None):
            super(MultiTaskLoss, self).__init__()
            self.emotion_weight = emotion_weight
            self.sentiment_weight = sentiment_weight
            self.emotion_class_weights = emotion_class_weights
            self.sentiment_class_weights = sentiment_class_weights

        def call(self, y_true, y_pred):
            # Separate emotion and sentiment losses
            emotion_loss = tf.keras.losses.sparse_categorical_crossentropy(
                y_true['emotion'], y_pred['emotion'], from_logits=True
            )
            sentiment_loss = tf.keras.losses.sparse_categorical_crossentropy(

```



```

# Compute softmax
y_pred_softmax = tf.nn.softmax(y_pred)

# Compute cross entropy
ce = -y_true_one_hot * tf.math.log(y_pred_softmax + 1e-7)

# Compute focal term
focal_term = tf.pow(1.0 - y_pred_softmax, self.gamma)

# Compute focal loss
focal_loss = focal_term * ce

# Apply alpha if provided
if self.alpha is not None:
    alpha_t = tf.gather(self.alpha, tf.cast(y_true, tf.int32))
    focal_loss = alpha_t * focal_loss

return tf.reduce_mean(tf.reduce_sum(focal_loss, axis=-1))

```

```

In [ ]: # Custom metrics
class MacroF1Score(tf.keras.metrics.Metric):
    """Macro F1 Score metric"""
    def __init__(self, num_classes, name='macro_f1', **kwargs):
        super(MacroF1Score, self).__init__(name=name, **kwargs)
        self.num_classes = num_classes
        self.precision = tf.keras.metrics.Precision()
        self.recall = tf.keras.metrics.Recall()

```

```

def update_state(self, y_true, y_pred, sample_weight=None):
    y_pred = tf.argmax(y_pred, axis=-1)
    self.precision.update_state(y_true, y_pred, sample_weight)
    self.recall.update_state(y_true, y_pred, sample_weight)

def result(self):
    precision = self.precision.result()
    recall = self.recall.result()
    f1 = 2 * (precision * recall) / (precision + recall + 1e-7)
    return f1

def reset_state(self):
    self.precision.reset_state()
    self.recall.reset_state()

```

5. Training Functions

```

In [ ]: def create_model(model_type, num_emotions, num_sentiments, tokenizer=None):
        """Factory function to create models based on type"""

        if model_type == 'lstm':
            if tokenizer is None:
                raise ValueError("LSTM model requires a tokenizer to get vocab_size")
            return LSTMEmotionModel(
                vocab_size=tokenizer.vocab_size,
                num_emotions=num_emotions,
                num_sentiments=num_sentiments

```

```

    )

elif model_type == 'bert':
    return BERTMultiTaskModel(
        model_name='bert-base-uncased',
        num_emotions=num_emotions,
        num_sentiments=num_sentiments
    )

elif model_type == 'roberta':
    return BERTMultiTaskModel(
        model_name='roberta-base',
        num_emotions=num_emotions,
        num_sentiments=num_sentiments
    )

elif model_type == 'dialoguernn':
    return DialogueRNN(
        base_model_name='bert-base-uncased',
        num_emotions=num_emotions,
        num_sentiments=num_sentiments
    )

elif model_type == 'cosmic':
    return COSMICModel(
        model_name='roberta-base',
        num_emotions=num_emotions,
        num_sentiments=num_sentiments
    )

```

```
else:
    raise ValueError(f"Unknown model type: {model_type}")
```

```
In [ ]: def create_callbacks(model_name, patience=5):
        """Create training callbacks"""
        callbacks = [
            ModelCheckpoint(
                f'best_{model_name}_model.weights.h5',
                monitor='val_loss',
                save_best_only=True,
                save_weights_only=True,
                verbose=1
            ),
            EarlyStopping(
                monitor='val_loss',
                patience=patience,
                restore_best_weights=True,
                verbose=1
            ),
            ReduceLROnPlateau(
                monitor='val_loss',
                factor=0.5,
                patience=patience//2,
                min_lr=1e-7,
                verbose=1
            ),
            tf.keras.callbacks.TensorBoard(
                log_dir=f'./logs/{model_name}',
```



```

        histogram_freq=1
    )
]
return callbacks

```

```

In [ ]: def train_model_simple(model, train_data, val_data, epochs=10, callbacks=None):
        """Train the model"""
        history = model.fit(
            train_data,
            validation_data=val_data,
            epochs=epochs,
            callbacks=callbacks,
            verbose=1
        )
        return history

def train_model(model_type, train_gen, val_gen, config, callbacks):
    """Train a specific model type"""

    # Create model
    tokenizer = train_gen.tokenizer
    num_emotions = len(json.load(open('data/label_mappings.json'))['emotion_to'])
    num_sentiments = len(json.load(open('data/label_mappings.json'))['sentiment_to'])

    model = create_model(
        model_type=model_type,
        num_emotions=num_emotions,
        num_sentiments=num_sentiments,
        tokenizer=tokenizer
    )

```

```

)

# Build model
dummy_input = next(iter(train_gen))[0]
_ = model(dummy_input)

# Compile model
optimizer = tf.keras.optimizers.Adam(learning_rate=config['learning_rate'])

model.compile(
    optimizer=optimizer,
    loss={
        'emotion': tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
        'sentiment': tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
    },
    metrics={
        'emotion': ['accuracy'],
        'sentiment': ['accuracy']
    },
    loss_weights={
        'emotion': config['emotion_weight'],
        'sentiment': config['sentiment_weight']
    }
)

# Train
history = model.fit(
    train_gen,
    validation_data=val_gen,

```

```

        epochs=config['num_epochs'],
        callbacks=callbacks,
        verbose=1
    )

    return model, history

```

```

In [ ]: import numpy as np
import tensorflow as tf
from sklearn.metrics import classification_report, confusion_matrix
from tqdm import tqdm

def evaluate_model(model, test_data, mappings):
    """Evaluate model and generate classification reports"""
    emotion_preds = []
    sentiment_preds = []
    emotion_labels = []
    sentiment_labels = []

    for batch in tqdm(test_data, desc='Evaluating'):
        inputs, labels = batch
        predictions = model(inputs, training=False)

        # Handle predictions (these are tensors)
        if isinstance(predictions['emotion'], tf.Tensor):
            emotion_preds.extend(tf.argmax(predictions['emotion'], axis=-1).numpy())
            sentiment_preds.extend(tf.argmax(predictions['sentiment'], axis=-1).numpy())
        else:
            emotion_preds.extend(np.argmax(predictions['emotion'], axis=-1))

```

```

        sentiment_preds.extend(np.argmax(predictions['sentiment'], axis=-1))

    # Handle labels (these might be numpy arrays or tensors)
    if isinstance(labels['emotion'], tf.Tensor):
        emotion_labels.extend(labels['emotion'].numpy())
        sentiment_labels.extend(labels['sentiment'].numpy())
    else:
        emotion_labels.extend(labels['emotion'])
        sentiment_labels.extend(labels['sentiment'])

# Convert to numpy arrays
emotion_preds = np.array(emotion_preds)
sentiment_preds = np.array(sentiment_preds)
emotion_labels = np.array(emotion_labels)
sentiment_labels = np.array(sentiment_labels)

# Generate reports
emotion_report = classification_report(
    emotion_labels, emotion_preds,
    target_names=list(mappings['idx_to_emotion'].values()),
    output_dict=True
)

sentiment_report = classification_report(
    sentiment_labels, sentiment_preds,
    target_names=list(mappings['idx_to_sentiment'].values()),
    output_dict=True
)

```

```
return {
    'emotion_report': emotion_report,
    'sentiment_report': sentiment_report,
    'emotion_preds': emotion_preds,
    'sentiment_preds': sentiment_preds,
    'emotion_labels': emotion_labels,
    'sentiment_labels': sentiment_labels
}
```

6. Training Pipeline

6.1 Train Pipeline - Baseline model

```
In [ ]: def main(model_type='bert'):
        # Configuration
        config = {
            'model_type': model_type, # 'bert', 'roberta', 'lstm', 'dialoguernn',
            'model_name': 'bert-base-uncased',
            'batch_size': 16,
            'learning_rate': 2e-5,
            'num_epochs': 10,
            'max_length': 128,
            'context_window': 3,
            'emotion_weight': 1.0,
            'sentiment_weight': 0.5
        }
```

```
# Load data
print("Loading data...")

train_df = pd.read_csv(train_path)
dev_df = pd.read_csv(dev_path)
test_df = pd.read_csv(test_path)

# Load label mappings
with open('data/label_mappings.json', 'r') as f:
    mappings = json.load(f)

# Prepare dialogues
train_dialogues, _, _ = prepare_for_transformers(train_df)
dev_dialogues, _, _ = prepare_for_transformers(dev_df)
test_dialogues, _, _ = prepare_for_transformers(test_df)

# Initialize tokenizer
tokenizer = AutoTokenizer.from_pretrained(config['model_name'])

# Set tokenizer properties
tokenizer.model_max_length = config['max_length']
if tokenizer.pad_token is None:
    tokenizer.pad_token = tokenizer.eos_token
    tokenizer.pad_token_id = tokenizer.eos_token_id

# Ensure padding is set correctly
tokenizer.padding_side = 'right'
```

```
# Create data generators
print("Creating data generators...")
train_gen = MELDDataGenerator(
    train_dialogues, tokenizer,
    batch_size=config['batch_size'],
    max_length=config['max_length'],
    context_window=config['context_window'],
    shuffle=True
)

val_gen = MELDDataGenerator(
    dev_dialogues, tokenizer,
    batch_size=config['batch_size'],
    max_length=config['max_length'],
    context_window=config['context_window'],
    shuffle=False
)

test_gen = MELDDataGenerator(
    test_dialogues, tokenizer,
    batch_size=config['batch_size'],
    max_length=config['max_length'],
    context_window=config['context_window'],
    shuffle=False
)

test_data_generator(train_dialogues[:5], tokenizer,
    batch_size=config['batch_size'],
    max_length=config['max_length'])
```

```

# Compute class weights
print("Computing class weights...")
emotion_labels = train_df['emotion_label'].values
sentiment_labels = train_df['sentiment_label'].values

emotion_weights = compute_class_weight(
    'balanced',
    classes=np.unique(emotion_labels),
    y=emotion_labels
)
sentiment_weights = compute_class_weight(
    'balanced',
    classes=np.unique(sentiment_labels),
    y=sentiment_labels
)

emotion_weights = tf.constant(emotion_weights, dtype=tf.float32)
sentiment_weights = tf.constant(sentiment_weights, dtype=tf.float32)

# Initialize model
print(f"Initializing {config['model_type']} model...")
if config['model_type'] == 'bert':
    model = BERTMultiTaskModel(
        model_name=config['model_name'],
        num_emotions=len(mappings['emotion_to_idx']),
        num_sentiments=len(mappings['sentiment_to_idx'])
    )
elif config['model_type'] == 'roberta':

```



```

        model = BERTMultiTaskModel(
            model_name='roberta-base',
            num_emotions=len(mappings['emotion_to_idx']),
            num_sentiments=len(mappings['sentiment_to_idx'])
        )
    elif config['model_type'] == 'lstm':
        # For LSTM, we need vocab size
        vocab_size = tokenizer.vocab_size
        model = LSTMEmotionModel(
            vocab_size=vocab_size,
            num_emotions=len(mappings['emotion_to_idx']),
            num_sentiments=len(mappings['sentiment_to_idx'])
        )
    elif config['model_type'] == 'cosmic':
        model = COSMICModel(
            model_name=config['model_name'],
            num_emotions=len(mappings['emotion_to_idx']),
            num_sentiments=len(mappings['sentiment_to_idx'])
        )

# Build model
dummy_input = next(iter(train_gen))[0]
_ = model(dummy_input)

# Compile model
loss = MultiTaskLoss(
    emotion_weight=config['emotion_weight'],
    sentiment_weight=config['sentiment_weight'],
    emotion_class_weights=emotion_weights,

```

```

        sentiment_class_weights=sentiment_weights
    )

    optimizer = tf.keras.optimizers.Adam(learning_rate=config['learning_rate'])

    model.compile(
        optimizer=optimizer,
        loss={
            'emotion': tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
            'sentiment': tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
        },
        metrics={
            'emotion': ['accuracy', MacroF1Score(len(mappings['emotion_to_idx']))],
            'sentiment': ['accuracy', MacroF1Score(len(mappings['sentiment_to_idx']))]
        },
        loss_weights={
            'emotion': config['emotion_weight'],
            'sentiment': config['sentiment_weight']
        }
    )

    # Print model summary
    model.summary()

    # Create callbacks
    callbacks = create_callbacks(config['model_type'])

    # Train model
    print("Starting training...")

```

```
history = train_model_simple(
    model,
    train_gen,
    val_gen,
    epochs=config['num_epochs'],
    callbacks=callbacks
)

# Evaluate on test set
print("\nEvaluating on test set...")
results = evaluate_model(model, test_gen, mappings)

print("\n=== Emotion Classification Report ===")
print(classification_report(
    results['emotion_labels'],
    results['emotion_preds'],
    target_names=list(mappings['idx_to_emotion'].values())
))

print("\n=== Sentiment Classification Report ===")
print(classification_report(
    results['sentiment_labels'],
    results['sentiment_preds'],
    target_names=list(mappings['idx_to_sentiment'].values())
))

# Plot results
plot_training_history(history)
plot_confusion_matrices(results, mappings)
```

```

    return model, history, results

def plot_training_history(history):
    """Plot training history"""
    fig, axes = plt.subplots(2, 2, figsize=(15, 10))

    # Emotion metrics
    axes[0, 0].plot(history.history['emotion_loss'], label='Train')
    axes[0, 0].plot(history.history['val_emotion_loss'], label='Val')
    axes[0, 0].set_title('Emotion Loss')
    axes[0, 0].set_xlabel('Epoch')
    axes[0, 0].legend()

    axes[0, 1].plot(history.history['emotion_accuracy'], label='Train')
    axes[0, 1].plot(history.history['val_emotion_accuracy'], label='Val')
    axes[0, 1].set_title('Emotion Accuracy')
    axes[0, 1].set_xlabel('Epoch')
    axes[0, 1].legend()

    # Sentiment metrics
    axes[1, 0].plot(history.history['sentiment_loss'], label='Train')
    axes[1, 0].plot(history.history['val_sentiment_loss'], label='Val')
    axes[1, 0].set_title('Sentiment Loss')
    axes[1, 0].set_xlabel('Epoch')
    axes[1, 0].legend()

    axes[1, 1].plot(history.history['sentiment_accuracy'], label='Train')
    axes[1, 1].plot(history.history['val_sentiment_accuracy'], label='Val')

```

```

axes[1, 1].set_title('Sentiment Accuracy')
axes[1, 1].set_xlabel('Epoch')
axes[1, 1].legend()

plt.tight_layout()
plt.savefig('training_history_tf.png')
plt.show()

def plot_confusion_matrices(results, mappings):
    """Plot confusion matrices"""
    fig, axes = plt.subplots(1, 2, figsize=(15, 6))

    # Emotion confusion matrix
    emotion_cm = confusion_matrix(
        results['emotion_labels'],
        results['emotion_preds']
    )
    sns.heatmap(
        emotion_cm, annot=True, fmt='d',
        xticklabels=list(mappings['idx_to_emotion'].values()),
        yticklabels=list(mappings['idx_to_emotion'].values()),
        ax=axes[0]
    )
    axes[0].set_title('Emotion Confusion Matrix')
    axes[0].set_xlabel('Predicted')
    axes[0].set_ylabel('True')

    # Sentiment confusion matrix
    sentiment_cm = confusion_matrix(

```

```

        results['sentiment_labels'],
        results['sentiment_preds']
    )
    sns.heatmap(
        sentiment_cm, annot=True, fmt='d',
        xticklabels=list(mappings['idx_to_sentiment'].values()),
        yticklabels=list(mappings['idx_to_sentiment'].values()),
        ax=axes[1]
    )
    axes[1].set_title('Sentiment Confusion Matrix')
    axes[1].set_xlabel('Predicted')
    axes[1].set_ylabel('True')

    plt.tight_layout()
    plt.savefig('confusion_matrices_tf.png')
    plt.show()

```

```

In [ ]: train_path = 'https://raw.githubusercontent.com/declare-lab/MELD/master/data/M
dev_path = 'https://raw.githubusercontent.com/declare-lab/MELD/master/data/MEI
test_path = 'https://raw.githubusercontent.com/declare-lab/MELD/master/data/ME
if __name__ == "__main__":
    model_types = ['bert', 'roberta', 'lstm', 'dialoguernn', 'cosmic']
    for model_type in model_types:
        model, history, results = main(model_type='bert')

```

Loading data...
Creating data generators...
Testing MELDDataGenerator...
Total batches: 4
Total samples: 58

Batch 0:

input_ids: (16, 128), dtype: int32
attention_mask: (16, 128), dtype: int32
emotion: (16,), dtype: int32
sentiment: (16,), dtype: int32

Batch 1:

input_ids: (16, 128), dtype: int32
attention_mask: (16, 128), dtype: int32
emotion: (16,), dtype: int32
sentiment: (16,), dtype: int32

Batch 2:

input_ids: (16, 128), dtype: int32
attention_mask: (16, 128), dtype: int32
emotion: (16,), dtype: int32
sentiment: (16,), dtype: int32

✓ Data generator test passed!
Computing class weights...
Initializing bert model...

Some weights of the PyTorch model were not used when initializing the TF 2.0 model TFBertModel: ['cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.bias', 'cls.seq_relationship.weight']

- This IS expected if you are initializing TFBertModel from a PyTorch model trained on another task or with another architecture (e.g. initializing a TFBertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing TFBertModel from a PyTorch model that you expect to be exactly identical (e.g. initializing a TFBertForSequenceClassification model from a BertForSequenceClassification model).

All the weights of TFBertModel were initialized from the PyTorch model.

If your task is similar to the task the model of the checkpoint was trained on, you can already use TFBertModel for predictions without further training.

Model: "bert_multi_task_model_6"

Layer (type)	Output Shape	Param
dropout_18 (Dropout)	?	
dense_24 (Dense)	(16, 768)	590,5
dropout_19 (Dropout)	?	
dense_25 (Dense)	(16, 384)	295,2
emotion (Dense)	(16, 7)	2,6
dense_26 (Dense)	(16, 768)	590,5
dropout_20 (Dropout)	?	
dense_27 (Dense)	(16, 384)	295,2
sentiment (Dense)	(16, 3)	1,1


Total params: 1,775,626 (6.77 MB)

Trainable params: 1,775,626 (6.77 MB)


Non-trainable params: 0 (0.00 B)

Starting training...


Epoch 1/10

625/625  **0s** 167ms/step - emotion_accuracy: 0.4327 - emotion_loss: 1.6219 - emotion_macro_f1: 0.9091 - loss: 2.1660 - sentiment_accuracy: 0.4307 - sentiment_loss: 1.0881 - sentiment_macro_f1: 0.7610


Epoch 1: val_loss improved from inf to 2.12456, saving model to best_bert_model.weights.h5

625/625  **162s** 196ms/step - emotion_accuracy: 0.4327 - emotion_loss: 1.6218 - emotion_macro_f1: 0.9091 - loss: 2.1659 - sentiment_accuracy: 0.4307 - sentiment_loss: 1.0881 - sentiment_macro_f1: 0.7610 - val_emotion_accuracy: 0.4265 - val_emotion_loss: 1.6159 - val_emotion_macro_f1: 0.9259 - val_loss: 2.1246 - val_sentiment_accuracy: 0.4454 - val_sentiment_loss: 1.0332 - val_sentiment_macro_f1: 0.7722 - learning_rate: 2.0000e-05


Epoch 2/10

625/625  **0s** 120ms/step - emotion_accuracy: 0.4626 - emotion_loss: 1.5432 - emotion_macro_f1: 0.9387 - loss: 2.0725 - sentiment_accuracy: 0.4655 - sentiment_loss: 1.0589 - sentiment_macro_f1: 0.7829


Epoch 2: val_loss improved from 2.12456 to 2.08970, saving model to best_bert_model.weights.h5

625/625  **83s** 133ms/step - emotion_accuracy: 0.4627 - emotion_loss: 1.5432 - emotion_macro_f1: 0.9388 - loss: 2.0725 - sentiment_accuracy: 0.4656 - sentiment_loss: 1.0588 - sentiment_macro_f1: 0.7830 - val_emotion_accuracy: 0.4265 - val_emotion_loss: 1.5937 - val_emotion_macro_f1: 0.9259 - val_loss: 2.0897 - val_sentiment_accuracy: 0.5059 - val_sentiment_loss: 1.0109 - val_sentiment_macro_f1: 0.7554 - learning_rate: 2.0000e-05


Epoch 3/10

625/625  **0s** 119ms/step - emotion_accuracy: 0.4763 - emotion_loss: 1.5095 - emotion_macro_f1: 0.9413 - loss: 2.0233 - sentiment_accuracy: 0.4975 - sentiment_loss: 1.0277 - sentiment_macro_f1: 0.7973


Epoch 3: val_loss improved from 2.08970 to 2.04334, saving model to best_bert_model.weights.h5

625/625  **83s** 132ms/step - emotion_accuracy: 0.4763 - emotion_loss: 1.5095 - emotion_macro_f1: 0.9413 - loss: 2.0233 - sentiment_accuracy: 0.4975 - sentiment_loss: 1.0277 - sentiment_macro_f1: 0.7973 - val_emotion_accuracy: 0.4292 - val_emotion_loss: 1.5539 - val_emotion_macro_f1: 0.9259 - val_loss: 2.0433 - val_sentiment_accuracy: 0.5383 - val_sentiment_loss: 0.9909 - val_sentiment_macro_f1: 0.7100 - learning_rate: 2.0000e-05


Epoch 4/10

625/625  **0s** 119ms/step - emotion_accuracy: 0.4773 - emotion_loss: 1.5049 - emotion_macro_f1: 0.9400 - loss: 2.0153 - sentiment_accuracy: 0.4960 - sentiment_loss: 1.0203 - sentiment_macro_f1: 0.7902


Epoch 4: val_loss improved from 2.04334 to 2.03082, saving model to best_bert_model.weights.h5

625/625  **82s** 132ms/step - emotion_accuracy: 0.4773 - emotion_loss: 1.5049 - emotion_macro_f1: 0.9400 - loss: 2.0153 - sentiment_accuracy: 0.4960 - sentiment_loss: 1.0203 - sentiment_macro_f1: 0.7902 - val_emotion_accuracy: 0.4256 - val_emotion_loss: 1.5482 - val_emotion_macro_f1: 0.9259 - val_loss: 2.0308 - val_sentiment_accuracy: 0.5374 - val_sentiment_loss: 0.9762 - val_sentiment_macro_f1: 0.7420 - learning_rate: 2.0000e-05

Epoch 5/10


625/625  **0s** 119ms/step - emotion_accuracy: 0.4814 - emotion_loss: 1.4860 - emotion_macro_f1: 0.9378 - loss: 1.9934 - sentiment_accuracy: 0.5050 - sentiment_loss: 1.0157 - sentiment_macro_f1: 0.7901

Epoch 5: val_loss improved from 2.03082 to 2.01919, saving model to best_bert_model.weights.h5


625/625  **82s** 132ms/step - emotion_accuracy: 0.4814 - emotion_loss: 1.4860 - emotion_macro_f1: 0.9378 - loss: 1.9934 - sentiment_accuracy: 0.5050 - sentiment_loss: 1.0157 - sentiment_macro_f1: 0.7901 - val_emotion_

accuracy: 0.4509 - val_emotion_loss: 1.5358 - val_emotion_macro_f1: 0.9264 - val_loss: 2.0192 - val_sentiment_accuracy: 0.5482 - val_sentiment_loss: 0.9761 - val_sentiment_macro_f1: 0.6935 - learning_rate: 2.0000e-05


Epoch 6/10

625/625  **0s** 123ms/step - emotion_accuracy: 0.4905 - emotion_loss: 1.4731 - emotion_macro_f1: 0.9437 - loss: 1.9729 - sentiment_accuracy: 0.5137 - sentiment_loss: 1.0001 - sentiment_macro_f1: 0.7912


Epoch 6: val_loss improved from 2.01919 to 1.99994, saving model to best_bert_model.weights.h5

625/625  **86s** 137ms/step - emotion_accuracy: 0.4905 - emotion_loss: 1.4732 - emotion_macro_f1: 0.9437 - loss: 1.9730 - sentiment_accuracy: 0.5137 - sentiment_loss: 1.0001 - sentiment_macro_f1: 0.7912 - val_emotion_accuracy: 0.4509 - val_emotion_loss: 1.5210 - val_emotion_macro_f1: 0.9241 - val_loss: 1.9999 - val_sentiment_accuracy: 0.5500 - val_sentiment_loss: 0.9640 - val_sentiment_macro_f1: 0.7046 - learning_rate: 2.0000e-05

Epoch 7/10

625/625  **0s** 129ms/step - emotion_accuracy: 0.4714 - emotion_loss: 1.4912 - emotion_macro_f1: 0.9352 - loss: 1.9906 - sentiment_accuracy: 0.5198 - sentiment_loss: 0.9991 - sentiment_macro_f1: 0.7828

Epoch 7: val_loss improved from 1.99994 to 1.99046, saving model to best_bert_model.weights.h5


625/625  **146s** 143ms/step - emotion_accuracy: 0.4714 - emotion_loss: 1.4912 - emotion_macro_f1: 0.9352 - loss: 1.9905 - sentiment_accuracy: 0.5198 - sentiment_loss: 0.9991 - sentiment_macro_f1: 0.7828 - val_emotion_accuracy: 0.4554 - val_emotion_loss: 1.5139 - val_emotion_macro_f1: 0.9197 - val_loss: 1.9905 - val_sentiment_accuracy: 0.5555 - val_sentiment_loss: 0.9613 - val_sentiment_macro_f1: 0.7548 - learning_rate: 2.0000e-05

Epoch 8/10


625/625  **0s** 122ms/step - emotion_accuracy: 0.4802 - emotio

n_loss: 1.4718 - emotion_macro_f1: 0.9360 - loss: 1.9682 - sentiment_accuracy: 0.5260 - sentiment_loss: 0.9928 - sentiment_macro_f1: 0.7863


Epoch 8: val_loss improved from 1.99046 to 1.98900, saving model to best_bert_model.weights.h5

625/625  **85s** 135ms/step - emotion_accuracy: 0.4802 - emotion_loss: 1.4718 - emotion_macro_f1: 0.9360 - loss: 1.9682 - sentiment_accuracy: 0.5260 - sentiment_loss: 0.9928 - sentiment_macro_f1: 0.7863 - val_emotion_accuracy: 0.4482 - val_emotion_loss: 1.5170 - val_emotion_macro_f1: 0.9223 - val_loss: 1.9890 - val_sentiment_accuracy: 0.5645 - val_sentiment_loss: 0.9523 - val_sentiment_macro_f1: 0.7481 - learning_rate: 2.0000e-05


Epoch 9/10

625/625  **0s** 125ms/step - emotion_accuracy: 0.4787 - emotion_loss: 1.4671 - emotion_macro_f1: 0.9355 - loss: 1.9618 - sentiment_accuracy: 0.5294 - sentiment_loss: 0.9893 - sentiment_macro_f1: 0.7909

Epoch 9: val_loss improved from 1.98900 to 1.96995, saving model to best_bert_model.weights.h5

625/625  **86s** 138ms/step - emotion_accuracy: 0.4788 - emotion_loss: 1.4671 - emotion_macro_f1: 0.9355 - loss: 1.9618 - sentiment_accuracy: 0.5294 - sentiment_loss: 0.9893 - sentiment_macro_f1: 0.7909 - val_emotion_accuracy: 0.4599 - val_emotion_loss: 1.5011 - val_emotion_macro_f1: 0.9185 - val_loss: 1.9699 - val_sentiment_accuracy: 0.5636 - val_sentiment_loss: 0.9440 - val_sentiment_macro_f1: 0.7251 - learning_rate: 2.0000e-05

Epoch 10/10

625/625  **0s** 119ms/step - emotion_accuracy: 0.4836 - emotion_loss: 1.4673 - emotion_macro_f1: 0.9349 - loss: 1.9600 - sentiment_accuracy: 0.5289 - sentiment_loss: 0.9864 - sentiment_macro_f1: 0.7893

Epoch 10: val_loss improved from 1.96995 to 1.95427, saving model to best_bert_model.weights.h5

625/625  **83s** 132ms/step - emotion_accuracy: 0.4836 - emotion_loss: 1.4673 - emotion_macro_f1: 0.9349 - loss: 1.9600 - sentiment_accuracy: 0.5289 - sentiment_loss: 0.9864 - sentiment_macro_f1: 0.7893

on_loss: 1.4673 - emotion_macro_f1: 0.9349 - loss: 1.9600 - sentiment_accuracy: 0.5289 - sentiment_loss: 0.9864 - sentiment_macro_f1: 0.7893 - val_emotion_accuracy: 0.4617 - val_emotion_loss: 1.4861 - val_emotion_macro_f1: 0.9062 - val_loss: 1.9543 - val_sentiment_accuracy: 0.5609 - val_sentiment_loss: 0.9429 - val_sentiment_macro_f1: 0.7223 - learning_rate: 2.0000e-05
Restoring model weights from the end of the best epoch: 10.

Evaluating on test set...

Evaluating: 100%|██████████| 164/164 [00:40<00:00, 4.06it/s]

```

=== Emotion Classification Report ===
      precision    recall  f1-score   support

    anger         0.52      0.19      0.27       345
    disgust        0.00      0.00      0.00        68
     fear         0.00      0.00      0.00        50
     joy          0.39      0.23      0.29       402
    neutral        0.54      0.95      0.69      1256
    sadness        0.00      0.00      0.00       208
    surprise       0.39      0.09      0.14       281

 accuracy         0.53         2610
 macro avg        0.26      0.21      0.20      2610
weighted avg        0.43      0.53      0.43      2610

```

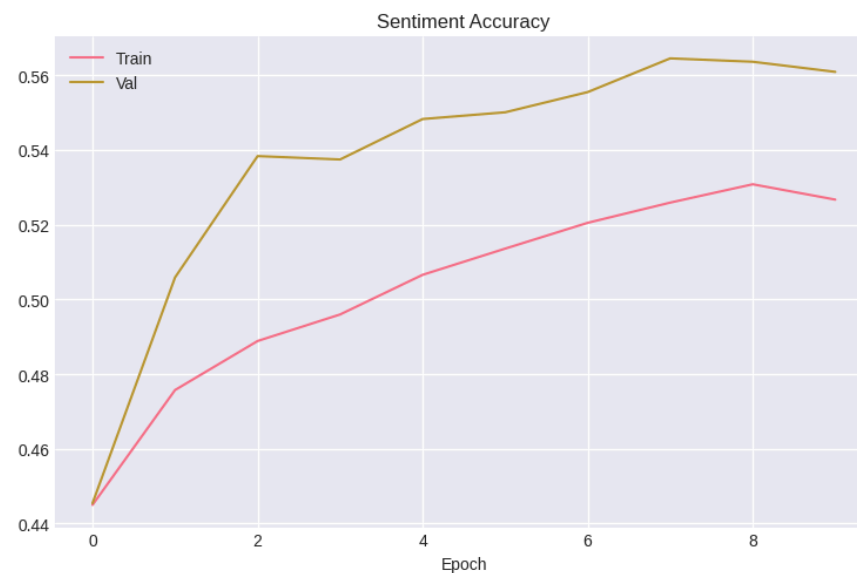
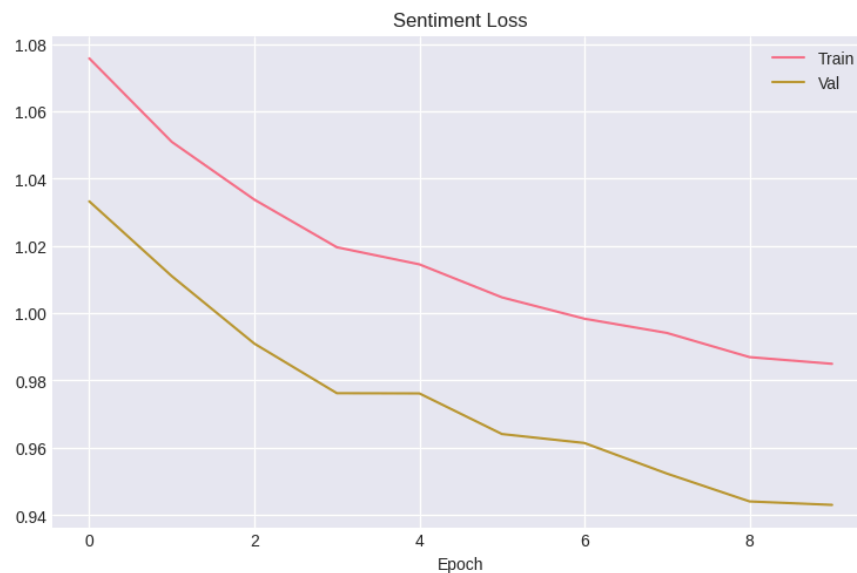
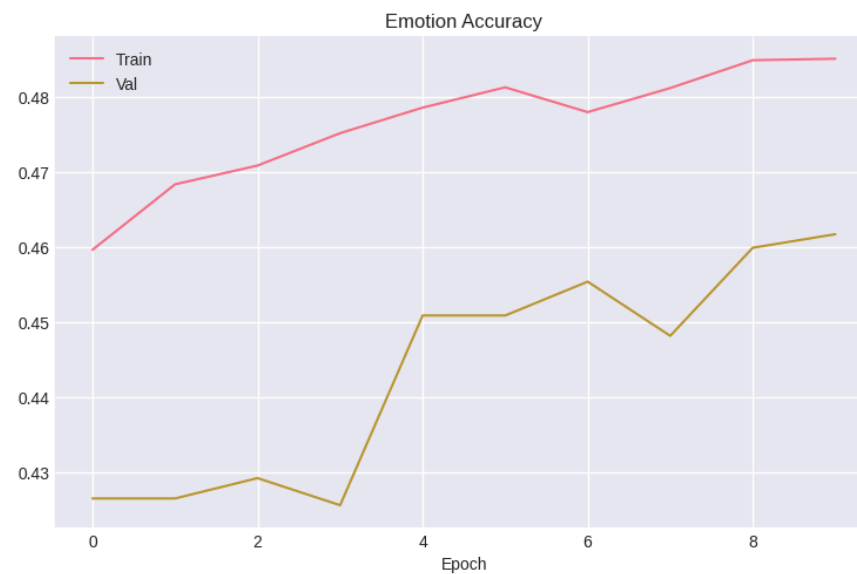
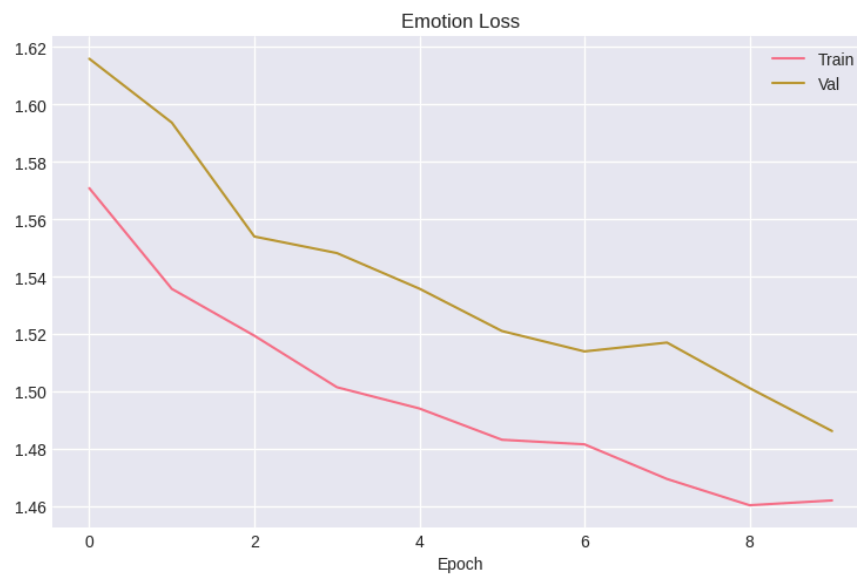
```


=== Sentiment Classification Report ===
      precision    recall  f1-score   support

 negative         0.52      0.63      0.57       833
   neutral        0.65      0.71      0.68      1256
  positive        0.53      0.24      0.33       521

 accuracy         0.59         2610
 macro avg        0.56      0.52      0.52      2610
weighted avg        0.58      0.59      0.57      2610

```





7. Inference and Deployment

```
In [ ]: class EmotionSentimentPredictor:
        """Inference class for emotion and sentiment prediction"""
        def __init__(self, model_path, model_type='bert', tokenizer_name='bert-base-uncased'):
            self.model_type = model_type
            self.tokenizer = AutoTokenizer.from_pretrained(tokenizer_name)

            # Load label mappings
            with open('data/label_mappings.json', 'r') as f:
                self.mappings = json.load(f)
```

```

# Initialize and load model
if model_type == 'bert':
    self.model = BERTMultiTaskModel(
        model_name=tokenizer_name,
        num_emotions=len(self.mappings['emotion_to_idx']),
        num_sentiments=len(self.mappings['sentiment_to_idx'])
    )

# Build model with proper dtype
dummy_input = {
    'input_ids': tf.zeros((1, 128), dtype=tf.int32),
    'attention_mask': tf.ones((1, 128), dtype=tf.int32)
}
_ = self.model(dummy_input)

# Load weights
self.model.load_weights(model_path)

@tf.function
def predict_batch(self, input_ids, attention_mask):
    """Optimized prediction for batches"""
    # Ensure proper dtypes
    input_ids = tf.cast(input_ids, tf.int32)
    attention_mask = tf.cast(attention_mask, tf.int32)

    inputs = {
        'input_ids': input_ids,
        'attention_mask': attention_mask
    }

```

```

    }

    predictions = self.model(inputs, training=False)

    emotion_probs = tf.nn.softmax(predictions['emotion'])
    sentiment_probs = tf.nn.softmax(predictions['sentiment'])

    emotion_preds = tf.argmax(emotion_probs, axis=-1)
    sentiment_preds = tf.argmax(sentiment_probs, axis=-1)

    return {
        'emotion_preds': emotion_preds,
        'sentiment_preds': sentiment_preds,
        'emotion_probs': emotion_probs,
        'sentiment_probs': sentiment_probs
    }

def predict(self, text, context=None):
    """Predict emotion and sentiment for single text"""
    # Prepare text
    if context:
        full_text = f"{context} [SEP] {text}"
    else:
        full_text = text

    # Tokenize
    encoded = self.tokenizer(
        full_text,
        padding='max_length',

```

```

        truncation=True,
        max_length=128,
        return_tensors='tf'
    )

    # Predict
    predictions = self.predict_batch(
        encoded['input_ids'],
        encoded['attention_mask']
    )

    emotion_idx = predictions['emotion_preds'][0].numpy()
    sentiment_idx = predictions['sentiment_preds'][0].numpy()

    return {
        'emotion': self.mappings['idx_to_emotion'][str(emotion_idx)],
        'sentiment': self.mappings['idx_to_sentiment'][str(sentiment_idx)],
        'emotion_probabilities': predictions['emotion_probs'][0].numpy().tolist(),
        'sentiment_probabilities': predictions['sentiment_probs'][0].numpy().tolist()
    }

def predict_dialogue(self, utterances, speakers=None):
    """Predict for entire dialogue"""
    predictions = []
    context = ""

    for i, utterance in enumerate(utterances):
        # Add speaker if available
        if speakers:

```

```

        utterance_with_speaker = f"{speakers[i]}: {utterance}"
    else:
        utterance_with_speaker = utterance

    # Predict
    pred = self.predict(utterance, context)
    predictions.append(pred)

    # Update context
    if context:
        context = f"{context} [SEP] {utterance_with_speaker}"
    else:
        context = utterance_with_speaker

    # Limit context length
    if len(context.split()) > 100:
        context = ' '.join(context.split()[-100:])

    return predictions

```

```

In [ ]: # Fixed TensorFlow Serving export function
def export_for_serving(model, export_path):
    """Export model for TensorFlow Serving with proper dtype handling"""

    # Create a serving function with explicit dtype conversions
    @tf.function
    def serving_fn(examples):
        # Parse the examples (assuming they come as strings)
        input_ids = examples['input_ids']

```

```

attention_mask = examples['attention_mask']

# Ensure proper dtypes
input_ids = tf.cast(input_ids, tf.int32)
attention_mask = tf.cast(attention_mask, tf.int32)

inputs = {
    'input_ids': input_ids,
    'attention_mask': attention_mask
}

predictions = model(inputs, training=False)

emotion_probs = tf.nn.softmax(predictions['emotion'])
sentiment_probs = tf.nn.softmax(predictions['sentiment'])

return {
    'emotion_logits': predictions['emotion'],
    'sentiment_logits': predictions['sentiment'],
    'emotion_probs': emotion_probs,
    'sentiment_probs': sentiment_probs,
    'emotion_predictions': tf.argmax(emotion_probs, axis=-1),
    'sentiment_predictions': tf.argmax(sentiment_probs, axis=-1)
}

# Get concrete function with proper signatures
concrete_serving_fn = serving_fn.get_concrete_function(
    examples={
        'input_ids': tf.TensorSpec(shape=[None, None], dtype=tf.int32, nar

```

```

        'attention_mask': tf.TensorSpec(shape=[None, None], dtype=tf.int32)
    }
)

# Save the model
tf.saved_model.save(
    model,
    export_path,
    signatures={'serving_default': concrete_serving_fn}
)
print(f"Model exported to {export_path}")

```

```

In [ ]: def export_keras_model(model, export_path):
        """Export model using Keras SavedModel format with proper file extension"""
        # Define the inference function
        @tf.function
        def inference(input_ids, attention_mask):
            # Ensure proper dtypes
            input_ids = tf.cast(input_ids, tf.int32)
            attention_mask = tf.cast(attention_mask, tf.int32)

            inputs = {
                'input_ids': input_ids,
                'attention_mask': attention_mask
            }

            return model(inputs, training=False)

        # Ensure directory exists

```

```

os.makedirs(os.path.dirname(export_path) if os.path.dirname(export_path)

# Add proper extension if not present
if not export_path.endswith('.keras') and not export_path.endswith('.h5'):
    export_path = export_path.rstrip('/') + '.keras'

# Save using Keras native format
model.save(export_path, overwrite=True)
print(f"Keras model saved to {export_path}")

```

```

In [ ]: # ===== Export Functions =====

def export_all_models():
    """Export all trained models for serving"""

    model_types = ['bert', 'roberta', 'lstm', 'dialoguernn', 'cosmic']

    for model_type in model_types:
        model_path = f'best_{model_type}_model.weights.h5'
        export_path = f'./saved_models/{model_type}/'

        try:
            print(f"\nExporting {model_type} model...")

            # Load model
            predictor = UnifiedEmotionSentimentPredictor(
                model_path=model_path,
                model_type=model_type
            )

```



```

        # Export
        predictor.model.save(export_path, save_traces=False)
        print(f"Exported to {export_path}")

    except Exception as e:
        print(f"Error exporting {model_type}: {e}")

```

```

In [ ]: def load_and_test_saved_model(saved_model_path):
        """Load and test the exported SavedModel"""
        # Load the SavedModel
        loaded_model = tf.saved_model.load(saved_model_path)

        # Get the serving function
        serving_fn = loaded_model.signatures['serving_default']

        # Test with dummy input
        test_input = {
            'input_ids': tf.constant([[101, 2054, 2003, 2115, 2171, 102, 0, 0]], dtype=tf.int32),
            'attention_mask': tf.constant([[1, 1, 1, 1, 1, 1, 0, 0]], dtype=tf.int32)
        }

        # Make prediction
        output = serving_fn(**test_input)

        print("SavedModel test successful!")
        print(f"Output keys: {list(output.keys())}")
        return output

```

```
In [ ]: def create_tflite_model(saved_model_path, tflite_path):
        """Convert SavedModel to TFLite for mobile/edge deployment"""
        # Load the converter
        converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_path)

        # Optimize for size (optional)
        converter.optimizations = [tf.lite.Optimize.DEFAULT]

        # Convert
        tflite_model = converter.convert()

        # Save
        with open(tflite_path, 'wb') as f:
            f.write(tflite_model)

        print(f"TFLite model saved to {tflite_path}")
```

```
In [ ]: # usage
        if __name__ == "__main__":
            try:
                # Initialize predictor
                predictor = EmotionSentimentPredictor(
                    model_path='best_bert_model.weights.h5', #best_bert_model.weights.h5
                    model_type='bert'
                )

                # Single prediction
                result = predictor.predict(
```

```

        text="I can't believe you did that!",
        context="We were supposed to meet at 5pm."
    )
print(f"Emotion: {result['emotion']}, Sentiment: {result['sentiment']}")

# Dialogue prediction
dialogue = [
    "How are you doing today?",
    "Not great, I lost my job.",
    "Oh no, I'm so sorry!",
    "Thanks, I'll be okay."
]

dialogue_predictions = predictor.predict_dialogue(dialogue)
for i, (utt, pred) in enumerate(zip(dialogue, dialogue_predictions)):
    print(f"Utterance: {utt}")
    print(f"    -> Emotion: {pred['emotion']}, Sentiment: {pred['sentiment']}")

# Single prediction
result = predictor.predict(
    text="I'm so frustrated with this service!",
    context="I've been waiting for 2 hours"
)

print(f"Emotion: {result['emotion']}, Sentiment: {result['sentiment']}")

#export_for_serving(model, './saved_model/1/')
# Export for serving – use the predictor's model
print("\nExporting model for TensorFlow Serving...")

```

```

export_for_serving(predictor.model, './saved_model/1/')

# Alternative: Export as Keras SavedModel
print("\nExporting as Keras SavedModel...")
export_keras_model(predictor.model, './keras_saved_model/')
except Exception as e:
    print(f"Error: {e}")
    print("\nTroubleshooting tips:")
    print("1. Ensure mixed precision is disabled for inference")
    print("2. Check that input tensors have correct dtypes (int32 for input_")
    print("3. Verify model weights are loaded correctly")
    print("4. Consider using export_keras_model() if export_for_serving() fa

```

Some weights of the PyTorch model were not used when initializing the TF 2.0 model TFBertModel: ['cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.bias', 'cls.seq_relationship.weight']

- This IS expected if you are initializing TFBertModel from a PyTorch model trained on another task or with another architecture (e.g. initializing a TFBertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing TFBertModel from a PyTorch model that you expect to be exactly identical (e.g. initializing a TFBertForSequenceClassification model from a BertForSequenceClassification model).

All the weights of TFBertModel were initialized from the PyTorch model.

If your task is similar to the task the model of the checkpoint was trained on, you can already use TFBertModel for predictions without further training.

Emotion: neutral, Sentiment: negative
Utterance: How are you doing today?
-> Emotion: neutral, Sentiment: neutral
Utterance: Not great, I lost my job.
-> Emotion: neutral, Sentiment: neutral
Utterance: Oh no, I'm so sorry!
-> Emotion: neutral, Sentiment: neutral
Utterance: Thanks, I'll be okay.
-> Emotion: neutral, Sentiment: neutral
Emotion: surprise, Sentiment: negative

Exporting model for TensorFlow Serving...
Model exported to ./saved_model/1/

Exporting as Keras SavedModel...
Keras model saved to ./keras_saved_model.keras

Summary

This TensorFlow implementation provides:

1. **Data Processing:** Custom data generators with batching and augmentation
2. **Model Architectures:**
 - LSTM with bidirectional layers
 - BERT/RoBERTa with TensorFlow integration
 - DialogueRNN with attention mechanisms

- COSMIC-style model with cross-attention

3. **Advanced Features:**

- Focal loss for extreme imbalance
- Text augmentation for minority classes
- Multimodal fusion architecture
- Mixed precision training support

4. **Training Pipeline:**

- Callbacks for checkpointing and early stopping
- TensorBoard integration
- Learning rate scheduling

5. **Production Features:**

- Optimized inference with `@tf.function`
- TensorFlow Serving export
- Batch prediction support

The implementation is modular and can be easily extended with new architectures or features.